

Classical baseline — MAE pretrain → finetune (paper recipe, AUC 0.97)

Stage 1 — self-supervised MAE pretrain (no_sub images only, no labels)

Stage 2 — supervised finetune (all 3 classes)

Scratch variant: random init instead of loading weights

